

Convolutional Neural Networks

ARISE 2021: ECE Machine Learning Lab

Haoran Zhu

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

July 27, 2021

Outline

- 1 Review of Neural Networks
- 2 Motivation
- 3 Dealing with Images in Computers
- 4 Convolution
- 5 Kernels

Extending Logistic Regression

- Motivation: Feature engineering in the model
 - Removes need for domain knowledge
 - Domain knowledge often doesn't exist: ex. object recognition
- Logistic Regression Model: $\hat{y} = \sigma(Wx + b)$
- Replace x with $z = f(Wx + b)$: $\hat{y} = \sigma(Wz + b)$
- So, $\hat{y} = \sigma(W_2 f(W_1 x + b_1) + b_2)$
- **Reminder:** all linear transforms can be represented as matrix multiplication
- We use non-linear function as f to give us a more expressive model
 - Recall polynomial transformations and exponential transformations of the data
 - These cannot be expressed as matrix multiplication

Extension to Neural Network


- Restrict $f(x)$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 x + b_1) + b_2)$
- We can optimize for both W_1, b_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$
 - Now we're *learning* the feature engineering
- But why stop here?...

Extension to Neural Network

- Restrict $f(x)$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 x + b_1) + b_2)$
- We can optimize for both W_1, b_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$
 - Now we're *learning* the feature engineering
- But why stop here?...

Extension to Neural Network

- Restrict $f(x)$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 x + b_1) + b_2)$
- We can optimize for both W_1, b_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$
 - Now we're *learning* the feature engineering
- But why stop here?...



Mathematical Model: Multi-Layer Perceptron


■ Model:

$$\hat{y} = f_{out}(W_{out}z_L + b_{out})$$

- Where, $z_l = f_l(W_l z_{l-1} + b_l)$ for $1 \leq l \leq L$, $z_0 = x$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform
- f_{out} depends on type of ML problem: (regression: linear, classification: sigmoid/soft-max)
 - **Regression:** Linear Output
 - **Binary Classification:** Sigmoid Output
 - **Multi-Class Classification:** Soft-max Output

Layers

- **Input:** feature vector, x
- **Output:** target vector, \hat{y}
 - linear/logistic regression
- **Hidden:** intermediate vectors, z or a
 - feature extraction



Common Activation Functions


- Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}$
 - $\sigma(z) \in (0, 1)$
- Tanh (hyperbolic tangent): $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$
 - $\tanh(z) \in [-1, 1]$
- ReLu (Rectified Linear Unit): $\text{relu}(z) = \max(0, z)$
 - easy to compute, performs well in practice

Guidelines for Designing a NN

- The design space for NN is HUGE
- Hyper-parameters so far:
 - L : # of layers
 - N_L : # hidden units per layer
 - f : activation function for each layer
 - bs : batch-size
 - lr : learning-rate
 - # of epochs
 - λ : weight-regularization constant
 - J : cost/loss function
- This can be overwhelming...

Guidelines for Designing a NN

- **Start Small:** 1 or 2 layers
 - # hidden units ~ 128
 - make sure code is working
 - increase size if val good
 - classification acc \geq guessing
- **One activation function**
 - for all hidden layers
- **Simple MLP Arch:**
 - Pyramid
 - Expand, combine & reduce



Outline

- 1 Review of Neural Networks
- 2 Motivation
- 3 Dealing with Images in Computers
- 4 Convolution
- 5 Kernels

Better performance with images

- Encoding locality
- How does an MLP see an image?
- Is this how we see images?

Examples: Lena & Mandrill




Outline

- 1 Review of Neural Networks
- 2 Motivation
- 3 Dealing with Images in Computers
- 4 Convolution
- 5 Kernels


Images in Computer

- Images are stored as arrays of quantized numbers in computers
- Gray scale image: 2D matrices with each entry specifying the intensity (brightness) of a pixel
- Pixel values range from 0 to 255, 0 being the darkest, 255 being the brightest



Color Images

- Color image: 3D array, 2 dimensions for space, 1 dimension for color
 - Can be thought of as three 2D matrices stacked together into a cube, each 2D matrix specifies the amount of each color: Red, Green, Blue value at each pixel



- Shape of this image: (1050,700,3)
- There are 1050x700 pixels, 3 channels: R,G,B

Outline

- 1 Review of Neural Networks
- 2 Motivation
- 3 Dealing with Images in Computers
- 4 Convolution
- 5 Kernels

Limitations of Fully Connected Network

- In Fashion MNIST, we used a fully connected network, in which each neuron in the hidden layer is connected to all $28 \times 28 = 784$ pixels
- Higher definition images often contain millions of pixels → It is not practical to use fully connected network
- Fully connected network treat each individual pixel as a feature, it does not utilize the positional relationship between pixels

Convolution

- Introducing a new operation: Convolution
- An operation on an image(matrix) X with a kernel W
- $Z = X \circledast W$

At each offset (j_1, j_2) compute:

$$Z[j_1, j_2] = \sum_{k_1=0}^{K_1-1} \sum_{k_2=0}^{K_2-1} W[k_1, k_2] X[j_1 + k_1, j_2 + k_2]$$

- Equation:

Example of a Convolution

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

Kernel

$$W = \begin{bmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{bmatrix}$$

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1


12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

Why Convolution?

- With convolution, each output pixel depends on only the neighboring pixels in the input
- This allows us to learn the positional relationship between pixels
- Use of different kernels allows us to detect features



Convolution for Multiple Channels

- A kernel for each channel. Could be same kernel, or different
- Perform a convolution for each of the channel, with the respective kernel
- Sum the results

Outline

- 1 Review of Neural Networks
- 2 Motivation
- 3 Dealing with Images in Computers
- 4 Convolution
- 5 Kernels**

Averaging Kernels

- Uniform Kernel: $\frac{1}{K_x K_y} \begin{bmatrix} 1 & .. & 1 \\ 1 & .. & 1 \\ 1 & .. & 1 \end{bmatrix}$


K_x = Number of Columns

K_y = Number of Rows

- Gaussian Kernel is a blurring kernel too.

Edge Detection

- Initial layers in a deep neural networks detect small patterns like lines, curves or edges.
- Subsequent layers combine these local features to create more complex features.



Edge Detection

- Using Sobel filters:

- Vertical Edge Detection $G_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$

- Horizontal Edge Detection $G_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Thank You!

- Next Class: Deep Learning and Applications of CNNs